

*We now understand the basic process lifecycle
& process control blocks.*


*Now, let us look into more
concrete details about these aspects & more!*

Process creation

- ◆ *Traditionally, OS transparently created all processes*
- ◆ Modern OSs provide *system call* for process creation
- ◆ First process (e.g., *init* in Unix) created at boot time
 - ◆ This process creates other processes at start up & later as needed
- ◆ Other processes can also create new processes as needed

Process creation steps

- ◆ Assign unique *identifier* to new process
- ◆ Allocate & set up memory space for process (*process memory image*)
 - ◆ Process control block (*PCB*)
 - ◆ Program & data → organized into regions
 - ◆ Code/text space
 - ◆ Global data space
 - ◆ Heap (dynamically allocated data) space
 - ◆ Stack (local function data) space
- ◆ Set up memory management structures for process
 - ◆ *We will look at details of these structures later...*
- ◆ Other structures OS may keep for performance monitoring, etc.



Process creation mechanisms


- ◆ 1st option → *cloning*
 - ◆ Process spawns process that is a copy of itself
 - ◆ Adopted by Unix based Oss → *fork()* system call
- ◆ 2nd option → *creation from scratch*
 - ◆ Process creates new process with appropriate parameters
 - ◆ Adopted by Windows → *CreateProcess()* system call

Unix process creation (*fork*)

- ◆ Process invokes *fork* to initiate creation of *new process*
- ◆ Creating process is *parent* & new process is *child*
- ◆ System call creates new process
- ◆ Data from parent process *copied* to memory of child process
 - ◆ Memory image, environment settings, I/O handles, etc.
 - ◆ Of course, *child* gets its own ID, scheduling info, etc.

17

If cloning is the only way to create processes, does that mean all processes can only execute identical code?



Not really!

- ◆ *fork* call returns a value
 - ◆ Upon success, fork returns child ID to parent process...
 - ◆ ...and returns 0 to child process
- ◆ Parent/child processes independently continue execution from statement after *fork*
- ◆ *[If process spawn fails, fork returns -1 to parent & no child is created]*

```
pid_t id = fork();  
if(id == -1) {  
    std::cout << "Error creating process\n";  
} else if (id == 0) {  
    std::cout << "I'm a child process!\n";  
} else {  
    std::cout << "I just became a parent!\n";  
}
```

In other words...

- ◆ ...return value to enable conditional execution after *fork*
 - ◆ So, even though program is same, paths can be different!
 - ◆ Different behavior can thus be achieved in parent & child

Alternatively

- ◆ Child can *change* the program it executes
 - ◆ Invoke *exec* family of system calls to change its memory image
 - ◆ Stops executing code in parent program & starts new program

```
pid_t id = fork();
if(id == -1) {
    std::cout << "Error creating process\n";
} else if (id == 0) {
    // child process functionality
    char* args[] = {"echo", "hello", NULL};
    execvp(args[0], args);
} else {
    std::cout << "I just became a parent!\n";
}
```


Must include <unistd.h>

Since cloning requires copying...

- ◆ ...it has high *overhead*
- ◆ Overhead somewhat reduced using *copy on write* concept
 - ◆ Start with both parent & child sharing same memory
 - ◆ Make copy if & when either process modifies data
 - ◆ So, if child changes memory image, copy need never be made!

Process termination

- ◆ As discussed earlier, process may terminate for many reasons
 - ◆ Voluntary exit upon task completion
 - ◆ Voluntary exit due to fatal error
 - ◆ Involuntary exit due to error/bug
 - ◆ Involuntary exit due to *kill* command by OS or other process
 - ◆ Of course, *killer* process must have appropriate *authorization*



On Unix based systems...

- ◆ OS maintains notion of process *hierarchy*
 - ◆ A process, its children, their children, etc. (i.e., all descendants) form process *group*
- ◆ Parent must be allowed to read child's *exit status*
- ◆ If *parent* terminates before *child*...
 - ◆ Child is now an *orphan* process
 - ◆ Orphan processes are *adopted* by *init* process

- ◆ If a *child* process terminates before *parent*
 - ◆ System will still need to keep child's *PCB*
 - ◆ Child process becomes a *zombie* process
 - ◆ “Dead”, but not “reaped”
 - ◆ Parent process can reap children by waiting for them to terminate
 - ◆ OS provides system call for this → *wait*



While processes are alive...

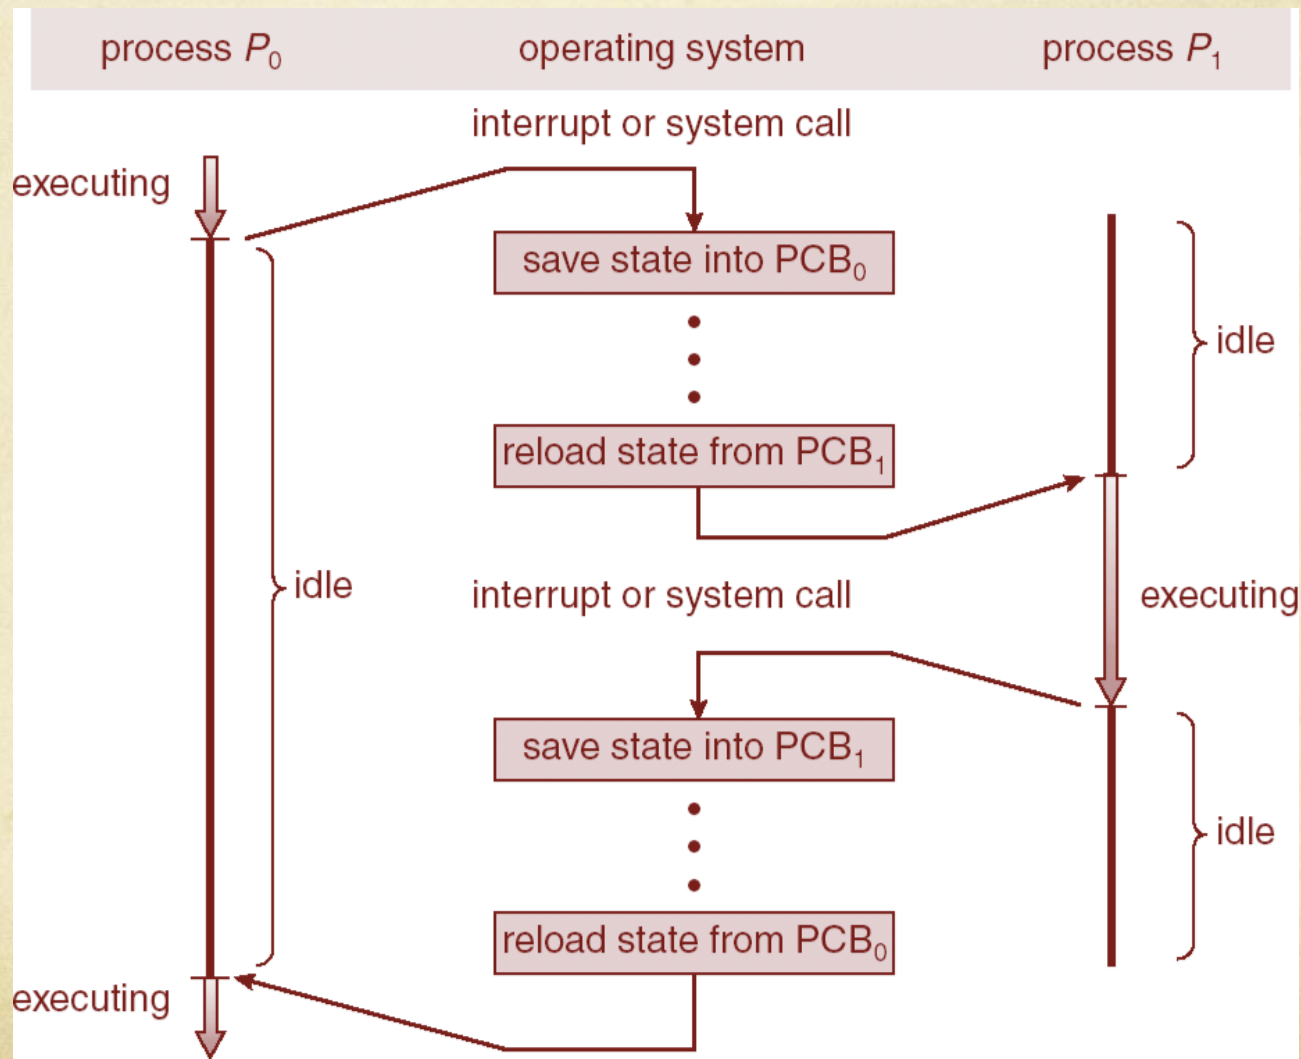
- ◆ ...they must be appropriately managed by OS
- ◆ Multiple processes may be concurrent
 - ◆ Big part of management is *switching* between processes
- ◆ OS needs to have access to PCBs of processes
 - ◆ Maintains *process table* to hold PCBs
- ◆ Questions that arise
 - ◆ *When* does OS switch to different process?
 - ◆ *Which* process does OS switch to?
 - ◆ *How* does OS perform process switch? —————→ *Mechanism*

} *Policies*

Process switching/context switching

- ◆ Example scenarios that may trigger process switch
 - ◆ Process termination (voluntary/involuntary)
 - ◆ New process activation
 - ◆ Executing process gets blocked (e.g., due to system call)
 - ◆ Event completion
 - ◆ Time slice expiration
- ◆ Process switching *mechanism* needs hardware & OS support
 - ◆ Requires transition to *kernel* mode → uses *interrupt*
 - ◆ Overhead of switch depends on hardware support (1-1000 μ s)

Process switch execution flow



Typical steps

- ◆ [*hw*] Save program counters & some registers on stack
- ◆ [*hw*] Load program counter specified in *interrupt vector*
 - ◆ Interrupt vector → has address of *interrupt service routine*
- ◆ [*asm*] Save context info (registers, etc.) in *PCB* of curr process
- ◆ [*asm*] Set up new stack
- ◆ [*C*] Remaining work for specific interrupt type
- ◆ [*C-sched*] Choose new process to be scheduled
- ◆ [*asm*] Restore context info for new process & start it