

Scheduling policies for interactive systems

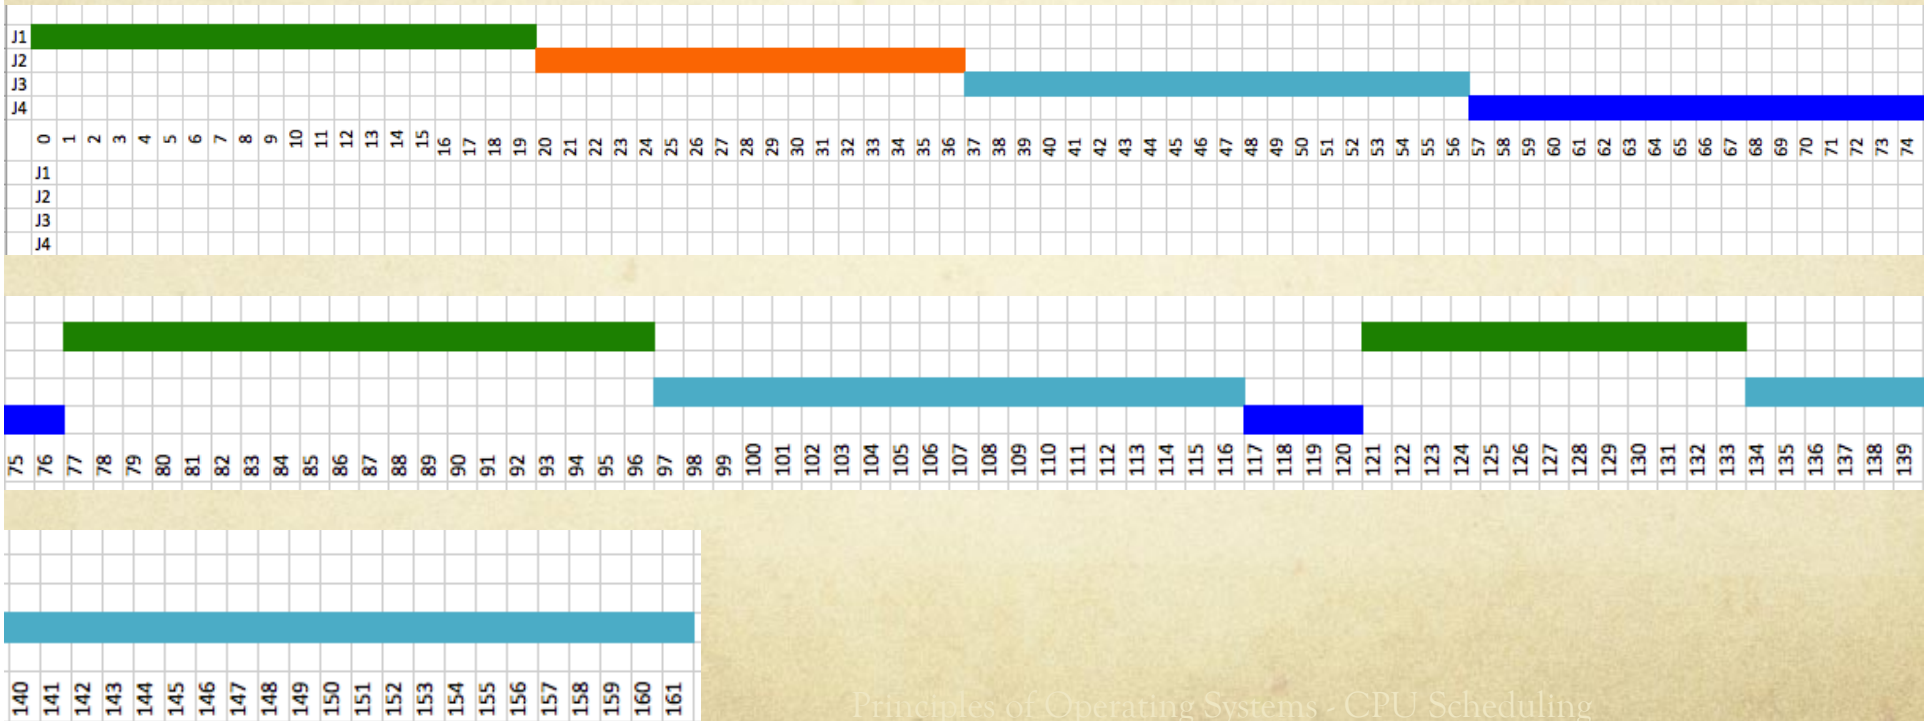
Round Robin (*RR*)

- ◆ Active jobs kept in *ready* queue
- ◆ Each job gets small unit (*quantum*) of CPU time
 - ◆ Quantum usually 10-100 milliseconds
- ◆ After quantum, job is *preempted* & next job in queue is scheduled
- ◆ *Pros*: simple; ensures fairness
- ◆ *Cons*: assumes all jobs are equally important
- ◆ Need to choose quantum *carefully*!
 - ◆ Too *short* → too much switching overhead
 - ◆ Too *long* → system not responsive

Round Robin Example

Quantum = 20

Job	CPU burst
J1	53
J2	17
J3	68
J4	24



Priority based scheduling

- ◆ *Priority* value (integer) associated with each job
- ◆ Job with *highest* priority is scheduled
- ◆ Schedule could be *preemptive* or *non-preemptive*
- ◆ Priority assignment based on *external* factor
 - ◆ Importance of job to user [typically *static* assignment]
 - ◆ *Cons*: low-priority jobs can *starve*
- ◆ Priority assignment based on *internal* factor
 - ◆ Aging
 - ◆ %CPU time used in last X hours

} [typically *dynamic* assignment]

Hybrid approaches may be used for mixed job types

Multilevel scheduling

- ◆ Ready queue *partitioned* into separate queues
 - ◆ E.g., system processes, foreground (*interactive*), background (*batch*)
- ◆ *Each* queue has its *own* scheduling algorithm
 - ◆ Example: foreground (*RR*), background(*FCFS*)
- ◆ Scheduling must be done between queues
 - ◆ *Fixed* priority – e.g, serve all foreground, then background
 - ◆ Possibility of starvation!
 - ◆ Time slice – e.g. 80% foreground (*RR*), 20% background (*FCFS*)