



Memory management

Context

- ◆ Process needs resources for execution
 - ◆ CPU cycles, I/O devices, etc.
 - ◆ *Memory* resources to store
 - ◆ Program code (“text”)
 - ◆ Data
 - ◆ OS structures for process

We will hereafter refer to the set of all process memory resources as process data

Context

- ◆ For process to execute, its data must be in *main memory*
 - ◆ Recall that main memory is *volatile*
- ◆ Process data stored in *non-volatile* memory, e.g., disk
- ◆ Process data must be *loaded* into main memory from disk

Simple scenario

- ◆ Allow only *one process* to be active at a time
 - ◆ *Load* process data from disk to main memory
 - ◆ Once process is done, *store* its contents to disk
- ◆ In reality, a process
 - ◆ May wait for I/O → poor *CPU utilization*
 - ◆ May not use all memory → poor *memory utilization*

More realistic scenario...

- ◆ Just like *CPU* multiplexed among multiple processes...
 - ◆ ...*memory* must be multiplexed among multiple processes
- ◆ What we need
 - ◆ *Multiple* processes should reside in memory *at a time*
 - ◆ Processes should not *collide* in memory
 - ◆ Process should not access other process' memory (*protection*)
 - ◆ Processes should be able to *share* memory if desired
- ◆ I.e., we need *memory protection* & *controlled overlap*

Memory partitioning


- ◆ Basic idea
 - ◆ Divide memory into multiple *partitions*
 - ◆ *Allocate* processes to partitions
 - ◆ Ensure *protection* across partitions
- ◆ Several questions arise
 - ◆ How should memory be divided into partitions?
 - ◆ Which process should be allocated to which partition?
 - ◆ Should partitions be allowed to change dynamically?
- ◆ Let's consider multiple options...

Static, equal-sized partitions

- ◆ *Statically* divide main memory into *equal-sized* partitions
- ◆ Map *each* process to *one* partition
 - ◆ For now, assume mapping does not change



Partitioned memory



Process memory footprints

Static, equal-sized partitions

- ◆ *Statically* divide main memory into *equal-sized* partitions
- ◆ Map *each* process to *one* partition
 - ◆ For now, assume mapping does not change



Partitioned memory



Process memory footprints

Static, equal-sized partitions

- ◆ *Statically* divide main memory into *equal-sized* partitions
- ◆ Map *each* process to *one* partition
 - ◆ For now, assume mapping does not change



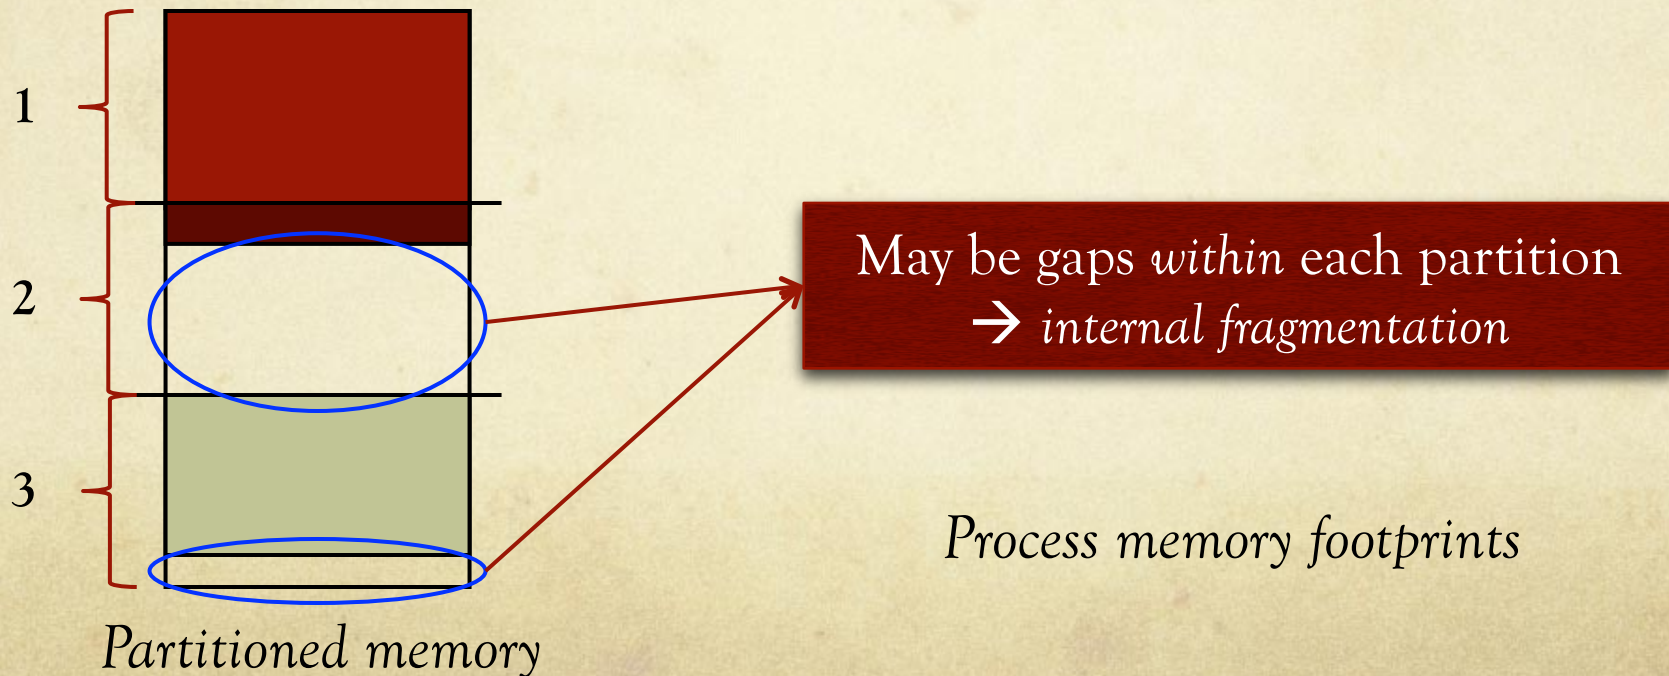
Partitioned memory



Process memory footprints


Static, equal-sized partitions

- ◆ *Statically* divide main memory into *equal-sized* partitions
- ◆ Map *each* process to *one* partition
 - ◆ For now, assume mapping does not change



Static, equal-sized partitions

- ◆ What if there are more processes than partitions?
 - ◆ Map multiple processes to *same* partition
 - ◆ *Store* existing process data to disk & *de-allocate* its partition
 - ◆ *Allocate* new process to partition & *load* its data from disk



Process memory footprints

Static, equal-sized partitions

- ◆ What if there are more processes than partitions?
 - ◆ Map multiple processes to *same* partition
 - ◆ *Store* existing process data to disk & *de-allocate* its partition
 - ◆ *Allocate* new process to partition & *load* its data from disk



Process memory footprints

Protection

- ◆ Ensure process only accesses locations in its partition
- ◆ Store start address or *base address* (*BA*) for each process
- ◆ Check every address issued by process
 - ◆ Must lie between *BA* & (*BA* + *partition size*)

Example

- ◆ Memory address range: 0 – 4999; Partitions of size 1000
- ◆ Partition to process mapping
 - ◆ Partition 1: P1, P6
 - ◆ Partition 2: P2, P7
 - ◆ Partition 3: P3
 - ◆ Partition 4: P4
 - ◆ Partition 5: P5
- ◆ Is P1 allowed to access address 1004?
 - ◆ No!
- ◆ Is P4 allowed to access address 3000?
 - ◆ Yes

Discussion

- ◆ *Pros*

- ◆ Very simple

- ◆ *Cons*

- ◆ Results in *internal* fragmentation
 - ◆ Results in *under-utilization* of memory
- ◆ Cannot fit process *larger* than partition size
- ◆ Takes a *one-size-fits-all* kind of approach